

Foundations of Data Science & Machine Learning

Summary — Week 02
Devansh Singh Rathore
111701011

B.Tech. in Computer Science & Engineering
Indian Institute of Technology Palakkad

March 14, 2021

Abstract

We study about Perceptron Learning Algorithm and its proof of upper bound of its computational complexity. Further we discuss non-linear separators, embedding in higher dimensional space using Kernel trick.

1 Perceptron learning algorithm

→ **Perceptron** is a supervised machine learning algorithm used for solving binary classification problems.

NOTE: We assume $b = 0$. Equivalently there exists a separating hyperplane passing through origin.

$$x = (x_1, x_2, \dots, x_n) \in \mathbb{R}^n, b \in \mathbb{R}$$

$$x' = (x_1, x_2, \dots, x_n, 1) \in \mathbb{R}^{(n+1)}$$

since $a = (a_1, a_2, \dots, a_n) \in \mathbb{R}^n$, assume $a' = (a_1, a_2, \dots, a_n, -b) \in \mathbb{R}^{(n+1)}$

For the good points we needed, $\langle x, a \rangle \geq b + \delta/2$

Now, $\langle x', a' \rangle = \langle x, a \rangle - b \geq \delta/2$

Similarly, for bad points we needed, $\langle x, a \rangle \leq b - \delta/2$ Now, $\langle x', a' \rangle = \langle x, a \rangle - b \leq -\delta/2$

Input: Two finite sets $G, B \subset \mathbb{R}^n$ which are linearly separable by a hyperplane passing through the origin.

Output: $a \in \mathbb{R}^n$ such that $\forall x \in G, \langle x, a \rangle > 0$ and $\forall x \in B, \langle x, a \rangle < 0$.

Algorithm: initially let's keep $a = 0$

Repeat until there is no update in a -

 for each $x \in G$:

 if($\langle x, a \rangle \leq 0$) :

 then $a = a + x$

 for each $x \in B$:

 if($\langle x, a \rangle \geq 0$) :

 then $a = a - x$

Claim: If G and B are linearly separable then the Perceptron algorithm terminates after $O(R^2/\delta^2)$ i.e. $O(1/\text{RelativeMargin}^2)$ updates, where

$$R = \max_{x \in G \cup B} \|x\|$$

$$\delta(\text{margin}) = \min_{x \in G, y \in B} \|x - y\|$$

$$\text{Relative Margin} = \delta/R$$

Proof: Since G and B are given as linearly separable. $\exists a^*$ s.t. $\forall x \in G, \langle x, a^* \rangle \geq \delta/2$ and $\forall x \in B, \langle x, a^* \rangle \leq -\delta/2$.

assume $\|a^*\| = 1$, (w.l.o.g.)

→ For $x \in G$, update occurs when $\langle x, a \rangle \leq 0$

$$a+ = x$$

$$\langle a, a^* \rangle + = \langle x, a^* \rangle$$

since $\langle x, a^* \rangle \geq \delta/2$, increase in $\langle a, a^* \rangle \geq \delta/2$

$\|a\|^2 = \|a + x\|^2$, so $\|a\|^2$ increases by $2 \langle x, a \rangle + \|x\|^2$ ($\langle x, a \rangle \leq 0$)

i.e. $\|a\|^2$ increases by $\leq R^2$

→ For $x \in B$, update occurs when $\langle x, a \rangle \geq 0$

$$a- = x$$

$$\langle a, a^* \rangle - = \langle x, a^* \rangle$$

since $-\langle x, a^* \rangle \geq \delta/2$, increase in $\langle a, a^* \rangle \geq \delta/2$

$\|a\|^2 = \|a - x\|^2$, so $\|a\|^2$ increases by $-2 \langle x, a \rangle + \|x\|^2$ ($\langle x, a \rangle \geq 0$)

i.e. $\|a\|^2$ increases by $\leq R^2$

→ After k updates,

$$\|a\| \geq \langle a, a^* \rangle \geq k \cdot \delta/2, \text{ so } \|a\| \geq k \cdot \delta/2 \quad \text{---(i)}$$

$$\text{while on the other hand } \|a\|^2 \leq k \cdot R^2 \text{ i.e. } \|a\| \leq \sqrt{k} \cdot R \quad \text{---(ii)}$$

From (i) and (ii), $k \leq 4 * R^2 / \delta^2$

Hence there cannot be more than $4 * R^2 / \delta^2$ updates in Perceptron algorithm.

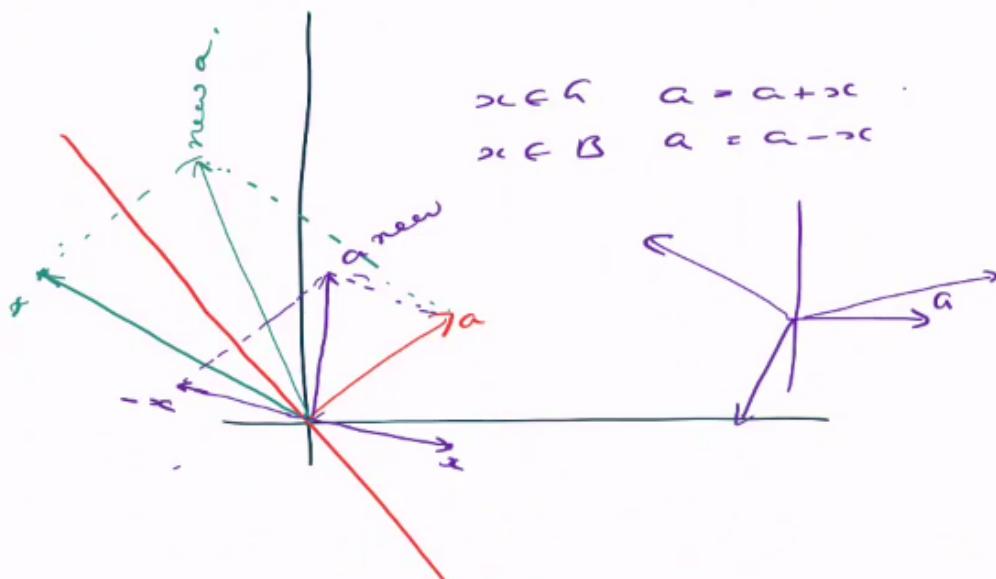


Fig 0.0 Working of Perceptron Algorithm

→ **Support Vector Machine (SVM)** is supervised machine learning algorithm used for optimal linear classification. SVM tries to find the maximum margin between Good and Bad points and finds suitable separating hyperplane.

2 Non-linear Separators using the Kernel Trick

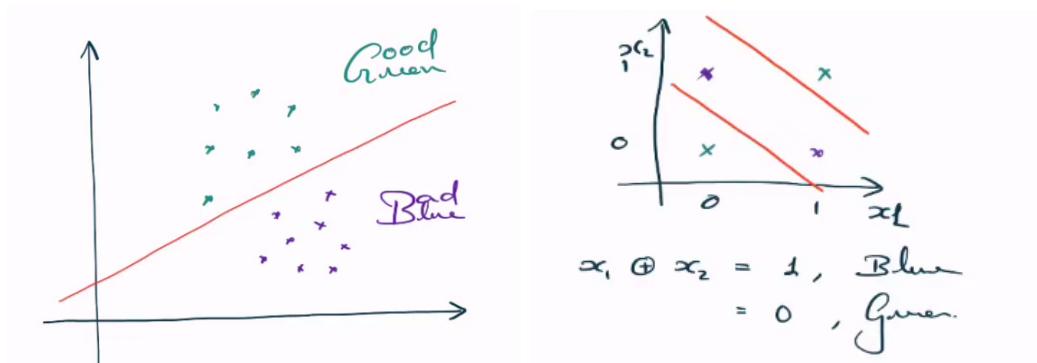


Fig 0.1 Linear Vs Non Linear Classification problem

→ We can have multiple linear classifier chunks for classifying points which cannot be classified using a single linear classifier hyperplane.

→ The base case of non-linear classification problem (Fig 0.1) is the XOR plot.

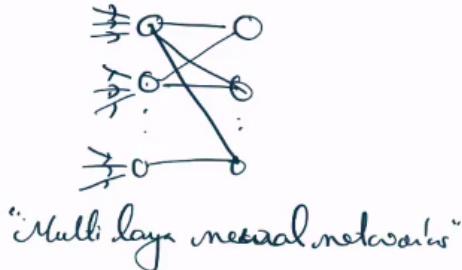


Fig 0.2 Multilayer Neural Network

2.1 Embedding in a higher dimensional space

→ Let's take an example where the Good points are within a circular boundary in a 2D plane and the Bad points lie outside the circular margin of radius R.

Now, the **Embedding function** is defined as: $\phi : \mathbb{R}^n \rightarrow \mathbb{R}^N$.

Here in this example, $\phi : \mathbb{R}^2 \rightarrow \mathbb{R}^3$

$$\phi((x_1, x_2)) = (x_1, x_2, x_1^2 + x_2^2)$$

→ Now in 3D space, the Good points lie below the plane $x_3 (= x_1^2 + x_2^2) = R^2$, while the Bad points lie above this plane. Hence we can say that the ϕ images of G and B are classified using **linear separator in \mathbb{R}^3** .

→ In general terms,

$$\phi : \mathbb{R} \rightarrow \mathbb{R}^{N=d+1}$$

$$\phi(x) = (1, x^1, x^2, \dots, x^d) \in \mathbb{R}^{d+1}$$

Separating hyperplane in \mathbb{R}^{d+1} looks like (a_0, a_1, \dots, a_d)

And for classification, we need to deal with polynomial equations like

$$a_0 + a_1 x^1 + a_2 x^2 + \dots + a_d x^d > 0$$

→ So you have at your disposal all polynomials of degree $\leq d$ as potential separators.

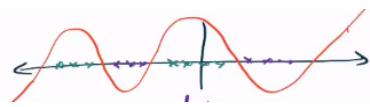


Fig 0.3 Polynomial Classifier

$\rightarrow \phi : \mathbb{R}^2 \rightarrow \mathbb{R}^N$

for d = 3, $\phi((x_1, x_2)) = (1, x_1, x_2, x_1^2, x_2^2, x_1x_2, x_1^3, x_2^3, x_1^2x_2, x_1x_2^2)$, so N = 10

$\rightarrow \phi : \mathbb{R}^3 \rightarrow \mathbb{R}^N$

for d = 2, n = 3, $\phi((x_1, x_2, x_3)) = (1, x_1, x_2, x_3, x_1^2, x_2^2, x_3^2, x_1x_2, x_1x_3, x_2x_3)$, so N = 10

\rightarrow In general, $N(n, d) = \text{number of terms with degree } \leq d \text{ in } (1 + x_1^1 + x_1^2 + \dots + x_1^d)(1 + x_2^1 + x_2^2 + \dots + x_2^d) \dots (1 + x_n^1 + x_n^2 + \dots + x_n^d)$

$$= C_d^{n+d}$$

= No. of monic monomials in n variables with degree $\leq d$.

\rightarrow Perceptron Learning algorithm with embedding

Input: $G = x_1, x_2, \dots, x_k \subseteq \mathbb{R}^n$

$B = x_{k+1}, x_{k+2}, \dots, x_l \subseteq \mathbb{R}^n$

$\phi : \mathbb{R}^n \rightarrow \mathbb{R}^N$

Output: $a \in \mathbb{R}^N, \forall x_j \in G, \langle a, \phi(x_j) \rangle > 0, \forall x_j \in B, \langle a, \phi(x_j) \rangle < 0$

Algorithm:

Initially $a = (0, 0, \dots) \in \mathbb{R}^N$

Repeat until there is no update in a:

 for each $x_j \in G$:

 if $\langle a, \phi(x_j) \rangle < 0, a = a + \phi(x_j)$

 for each $x_j \in B$:

 if $\langle a, \phi(x_j) \rangle > 0, a = a - \phi(x_j)$

\rightarrow This will terminate when $\phi_G = \{\phi(x) : x \in G\}$ & $\phi_B = \{\phi(x) : x \in B\}$ are linearly separable in \mathbb{R}^N .

\rightarrow This algorithm is computationally expensive when N is very large. We can make it efficient using the "Kernel trick".

2.2 Kernel Trick

\rightarrow Observation 1: Most of the computations are inner products of the form $\langle a, \phi(x) \rangle$.

\rightarrow Observation 2: a is a linear (integer) combination of $\phi(x_i), i \in G \cup B$ i.e.

$$a = \sum_{j=1}^l \alpha_j \phi(x_j)$$

where $\alpha_j = \text{isBad}(x_j) * \text{number of times } x_j \text{ passed the update condition}$

$\text{isBad}(x_j) = -1$ if $x_j \in B$, else 1.

\rightarrow From Observation 1 & 2 -

$$\langle a, \phi(x_i) \rangle = \left\langle \sum_{j=1}^l \alpha_j \phi(x_j), \phi(x_i) \right\rangle$$

$$= \sum_{j=1}^l \alpha_j \langle \phi(x_j), \phi(x_i) \rangle, \text{ where } x_i, x_j \in G \cup B$$

\rightarrow So the problem of computational efficiency boils down to checking whether we can compute $\langle \phi(x_i), \phi(x_j) \rangle$ efficiently i.e. with complexity/time independent of N.

if $\phi : \mathbb{R} \rightarrow \mathbb{R}^{d+1}$

$$\phi(x) = (1, x, x^2, \dots, x^d)$$

$$\phi(y) = (1, y, y^2, \dots, y^d)$$

$$\langle \phi(x), \phi(y) \rangle = 1 + xy + x^2y^2 + \dots + x^dy^d$$

$$= (1 - (xy)^{d+1}) / (1 - (xy)), \text{ using G.P.}$$

\rightarrow Perceptron Learning algorithm with Kernel Trick

Input: $G = x_1, x_2, \dots, x_k \subseteq \mathbb{R}^n$

$B = x_{k+1}, x_{k+2}, \dots, x_l \subseteq \mathbb{R}^n$

$K : \mathbb{R}^n \times \mathbb{R}^n \rightarrow \mathbb{R}$, $K(x, y) = \langle \phi(x), \phi(y) \rangle$

Output: $(\alpha_1, \alpha_2, \dots, \alpha_l) \in \mathbb{Z}^l$, $\forall x_i \in G, \sum_{j=1}^l \alpha_j K(x_j, x_i) > 0$, $\forall x_i \in B, \sum_{j=1}^l \alpha_j K(x_j, x_i) < 0$

Algorithm:

Initially $\alpha = (0, 0, \dots) \in \mathbb{R}^N$

Repeat until there is no update in α :

 for each $x_j \in G$:

 if $\sum_{i=1}^l \alpha_i K(x_j, x_i) \leq 0$, $\alpha_j + = 1$

 for each $x_j \in B$:

 if $\sum_{i=1}^l \alpha_i K(x_j, x_i) \geq 0$, $\alpha_j - = 1$